

Ghost Story Lens

Design Rationale and Build Intent

Prepared by Jong-Keyong (Jake) Kim, with Claude · 19 July 2026

Purpose

This report sets down, in one place, how I wanted Ghost Story Lens built and why, as it developed across a series of working sessions on the Product Requirements Document. It is meant as a companion to the PRD rather than a replacement for it. The PRD tells a builder what to make. This report explains the reasoning that produced those requirements, so that a future reader, a committee member, or an agent picking the project back up mid-stream can see the argument behind the specification and not just its output.

Starting Point

The project began from a fully drafted v1.0 PRD: a genre-fit analyzer that scores a submitted story on a 1–10 Ghost-Story Fit spectrum against a versioned, revisable rubric, and visualizes the haunting through a Relational Manifestation Map built on a ghost, mediator, and seer model drawn from my own dissertation research. That starting document already carried the premise I most wanted to protect through every later revision. Genre classification is an interpretive argument, not an objective discovery. Every change documented below was tested against whether it strengthened or diluted that premise.

Visual Identity

The first substantive revision was a full rejection of an earlier visual design pass (produced by a different model) that I felt had drifted into generic, uninspired styling. Rather than hand over a mood board, I asked for a design built from the product's own argument.

The governing decision was to stage a visible tension between two registers rather than illustrate the genre directly. An analytical register, carried by navy and graphite structure, grids, and legible tables, represents the AI's attempt to impose rubric-based order on a genre that resists it. A spectral register, expressed almost entirely through a single signature interaction, a soft point of light that trails and fades behind the user's cursor like a will-o'-the-wisp, represents the part of the genre that stays ungraspable no matter how much structure is layered on top of it. The interface argues its own premise instead of merely describing it.

Alongside that, I set a hard rule that has held through every later revision: no literal Gothic or horror iconography anywhere in the product. No skulls, candles, fog overlays, ruined castles, or ghost silhouettes. When I later asked for a Gothic-flavored icon for the theme toggle, this rule was still in force, so the request was resolved as an architectural reference (a Gothic lancet arch, the pointed window

silhouette of a cathedral) rather than a horror motif. Gothic as a formal, art-historical category was welcome. Gothic as horror-genre shorthand was not.

Readability and Layout

I asked for two corrections to how text behaves on the page. First, the reading column needed to switch from a fixed pixel width to a fluid width proportional to the viewport, because a fixed column on a wide screen was producing unwanted whitespace and premature line breaks. Second, after an intermediate attempt to enlarge body text to improve readability, I reverted that decision and fixed body text at 13 pt across all breakpoints. Readability was the goal in both cases; a larger fixed size was not the correct means to it.

Information Architecture

The navigation was restructured around a deliberately quiet entry point. The Home page carries almost no text, just a single inviting icon (the will-o'-the-wisp motif doubling as the entry action) that leads the visitor into an About page before anything else. From there the order runs About, Sample Analysis, Analyze, My Analysis. The About page absorbed expanded usage guidance for the three features most likely to confuse a first-time user: multilingual output, the interactive map, and the AI research conversation, each explained in plain terms rather than left to be discovered by trial and error.

My Analysis was added as a new page in its own right, not an afterthought bolted onto Analyze. It carries the full range I wanted a returning user to have: a list of saved analyses, one-click reopening, rerun and rubric-version comparison, deletion, and re-access to prior exports without regenerating them.

Voice and Vocabulary

Once the shape of the product was settled, I asked for a pass across the entire PRD, not just the user-facing copy, to recalibrate vocabulary for the actual audience, namely undergraduates, graduate students, creative writers, and English department faculty. Terms that were serving no purpose beyond signaling technical fluency, deterministic, structured output, schema, idempotent, adversarial, were rewritten in plainer language wherever they appeared in explanatory prose. I drew one boundary deliberately. The implementation-facing architecture section was left alone, since a Replit developer still needs that vocabulary to build the thing, and diluting it there would have cost precision for no readership benefit. The rule was audience-appropriate language, not uniform simplicity.

Export Behavior

The export model was simplified from three separate downloads (an analysis-summary JPG, a network-map JPG, and a combined PDF) to a single combined PDF, and the action itself was renamed from Download Results to Download Result to match. A user should not have to decide which of three files they want; the combined report already contains everything the two JPGs would have shown separately.

Small Points of Care

A few details were corrected because they affect whether the product feels considered rather than assembled. The creator's contact email on the About page was made a functional mailto link rather than inert text, so a question or correction is one click away from being sent. These are minor individually, but I wanted them treated with the same seriousness as the larger architectural decisions, since a reader's trust in the tool is built out of details like this as much as out of the rubric's rigor.

Light and Dark Mode

A theme toggle was added to the top right of the primary navigation, on every page, expressed as a Gothic lancet-arch icon holding a sun disc in Light Mode and a crescent moon in Dark Mode. The arch itself never changes shape, so it functions as a stable brand mark rather than a decoration that redraws itself. Dark Mode is a re-tuned version of the same five-role palette rather than a separate design, which keeps the analytical/spectral duality intact in both states. One exception was set deliberately: the combined PDF export always renders in the Light Mode palette regardless of the user's on-screen theme, so a printed or archived report stays legible independent of how any one reader has their screen configured.

Feature Scope: Custom Rubric Mode

When the Replit build agent flagged Custom Rubric Mode as a Should-priority, open-decision item and proposed deferring it, I used that flag as the occasion to resolve the ambiguity rather than simply wave the agent past it. I committed the feature to this build, raised its priority to Must, and specified it properly. A signed-in account gate, editable criteria with name, definition, and weight fields, a hard validation rule that weights must sum to 100 percent, editable score-band boundaries constrained to stay monotonic and non-overlapping across the full 1.0-to-10.0 spectrum, and optional exclusion rules that can cap or floor a score outside the normal weighted calculation. Critically, a user's edited rubric follows the same proposal, approval, versioning, and audit trail already built for the AI's own rubric-gap suggestions, so a human-authored change and a model-authored change are governed identically once submitted.

This exchange captures something about how I want to work with an agent through this build. A flag raised by the agent is not a request for permission to skip something; it is an invitation for me to make the call explicitly and put the answer back into the specification, so the next flag on the same topic never has to be raised again.

Instruction Given to the Build Agent

The build was kicked off with a short, direct brief rather than a long one. Build to the PRD and the mock-ups exactly, treat the PRD as the source of truth for requirements and the mock-ups as visual reference rather than pixel templates, start with the Corpus-Calibrated pipeline before touching the other

two analytical modes, hold the line on the non-negotiables (no literal Gothic imagery, fit and confidence and entity type always reported separately, every evidentiary claim traceable to a stored passage, deterministic application logic computing final scores rather than the model itself, and the PDF export always rendered in Light Mode), and flag anything ambiguous rather than improvise past it.

Underlying Principles

Read together, these decisions rest on a small number of recurring commitments that I want any future collaborator on this project, human or model, to hold onto even where this report has not spelled out a specific instance:

- The interface should argue the product's premise, not merely decorate it. Design choices are accountable to the same intellectual claim the rubric makes.
- Symbolic evocation of the genre is welcome; literal genre iconography is not. This applies to icons, copy, and interaction design alike.
- Vocabulary should match the reader in front of it. A dissertation committee member and a Replit developer are not the same reader, and the same document can serve both without flattening to the lowest common denominator.
- A tool that revises its own rubric needs the same rigor applied to human-authored revisions as to AI-proposed ones. Neither kind of change should be able to quietly overwrite the default.
- Ambiguity flagged during the build is an opportunity to decide, not a reason to defer indefinitely. Once decided, the decision belongs in the PRD, not only in conversation.

Status at Time of Writing

The PRD stands at version 1.4. Six page mock-ups (Home, About, Sample Analysis, Analyze, Results, My Analysis) exist in Day Mode reflecting the current specification, including the theme-toggle icon and the single combined-PDF export action. Custom Rubric Mode is now fully specified and committed to the MVP. The Replit build agent has been instructed to begin with the Corpus-Calibrated pipeline and to flag rather than improvise past any remaining ambiguity.